

NL2SQL USING LARGE LANGUAGE MODEL

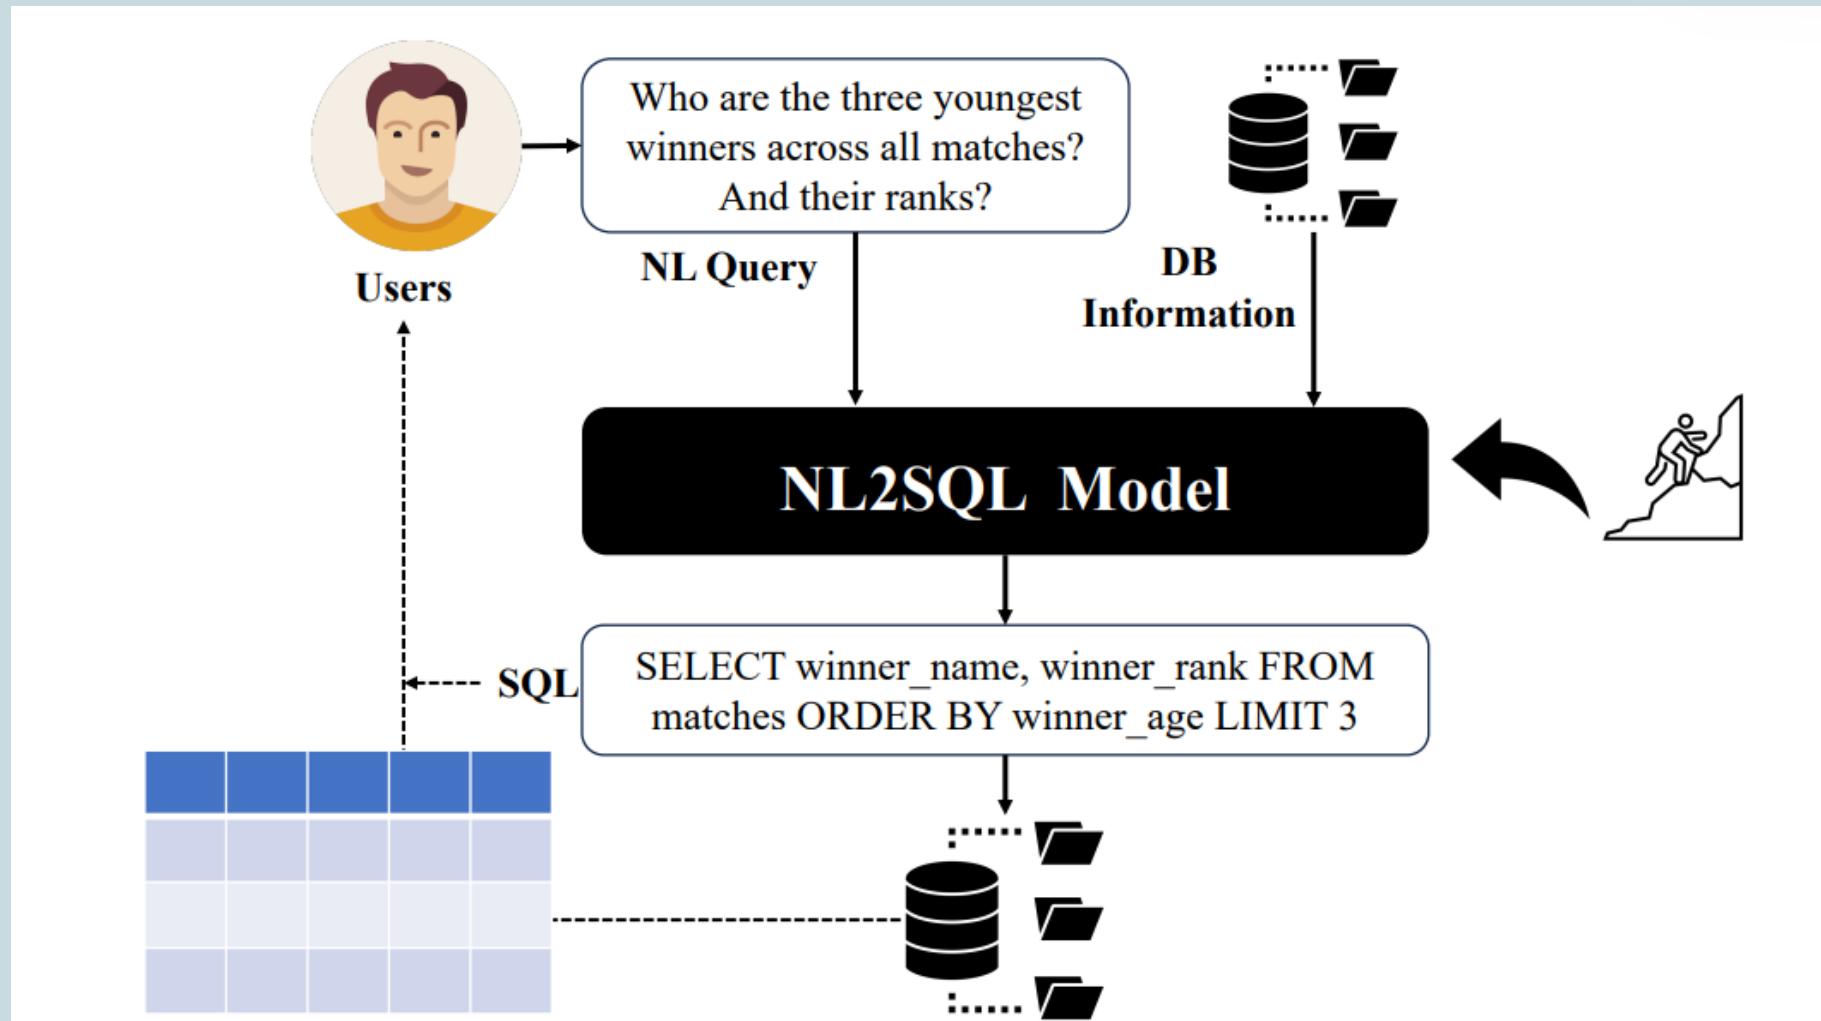


Anubhav Binit

OUTLINE

- NL2SQL PROBLEM AND BACKGROUND PRODUCT
- LANGUAGE MODEL-POWERED NL2SQL SOLUTIONS
 - Pre-Processing
 - Schema Linking
 - Database Content Retrieval
 - NL2SQL Translation Methods
 - Encoding Strategy
 - Decoding Strategy
 - Intermediate Representation
 - Post-Processing
- EVALUATION METRIC

WHAT IS NATURAL LANGUAGE TO SQL (NL2SQL)?



NL2SQL TASK CHALLENGES

C1. Ambiguous NL Query

Natural language queries can be vague or have multiple interpretations.

Example: *"List the top customers."* → Top by what? Revenue, number of purchases, or something else?

C2. Database Schema Dependency

The correct SQL depends on the database structure, which varies by case.

Example: Same NL query may need different SQL if column or table names differ across databases.

C3. Complex Schema

Large or deeply relational schemas make it hard to find the right tables and columns.

Example: A schema with 50+ tables and foreign keys is hard to navigate

C4. Multiple Possible SQL Queries

Different SQL queries can give valid answers to the same NL query.

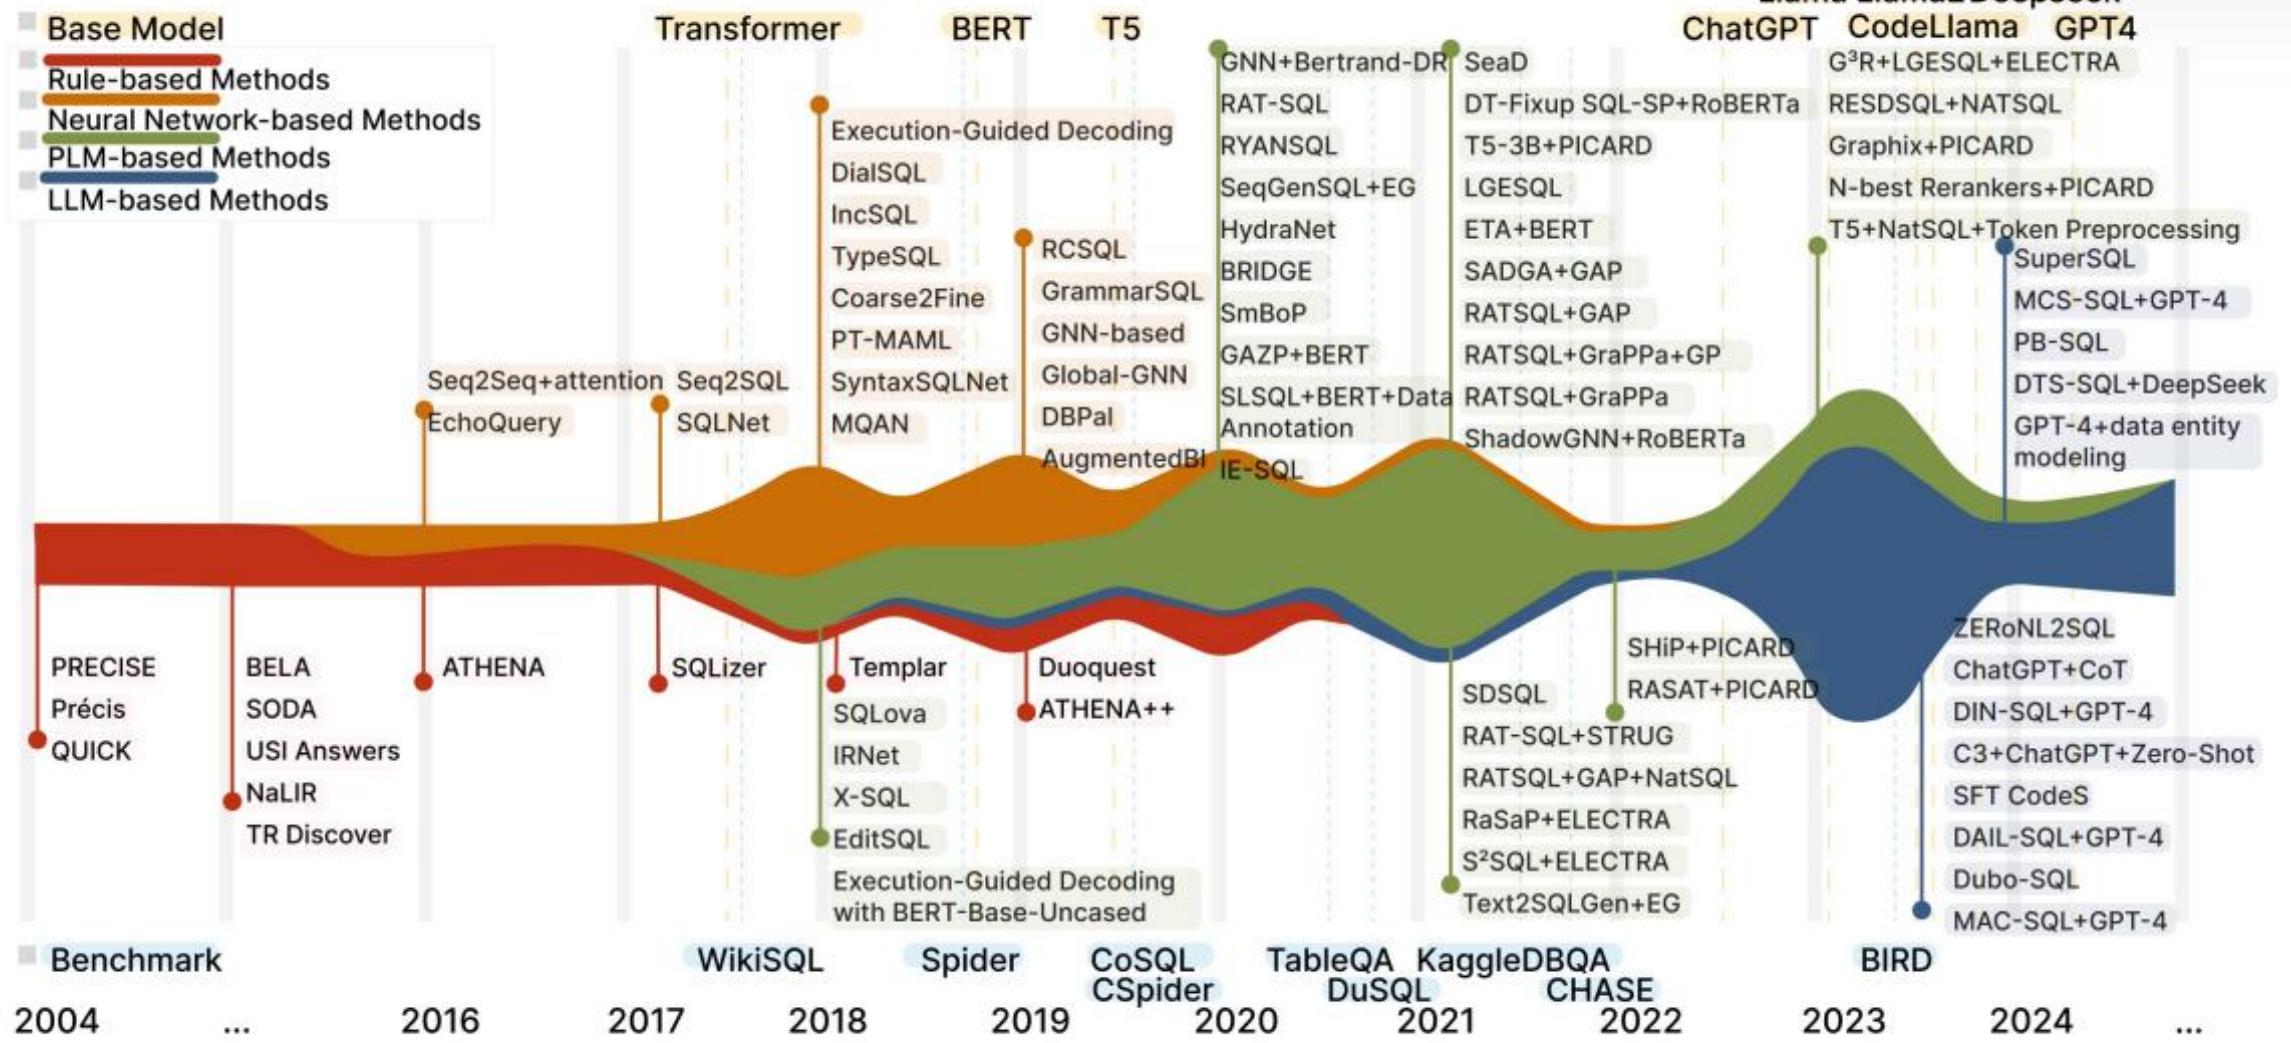
Example: *"Show all employees in the Sales department."* → Can use JOINs or subqueries.

C5. Database Domain Adaptation

A model trained on one domain (e.g., movies) may not perform well on another (e.g., healthcare).

Example: Column names and query patterns change between domains.

NL2SQL Model Evolution Stream Graph



Liu, X., Shen, S., Li, B., Ma, P., Jiang, R., Zhang, Y., Fan, J., Li, G., Tang, N., & Luo, Y. (2024). *A Survey of NL2SQL with Large Language Models: Where are we, and where are we going?*

AN OVERVIEW OF LANGUAGE MODEL-POWERED NL2SQL

- Pre-Processing
 - Schema Linking
 - Database Content Retrieval
 - Additional Information Acquisition
- NL2SQL Translation Methods
 - Encoding Strategy
 - Decoding Strategy
 - Task-specific Prompt Strategies
 - Intermediate Representation
- Post-Processing
 - Correction
 - Consistency
 - Execution- Guided
 - N-best Rerankers

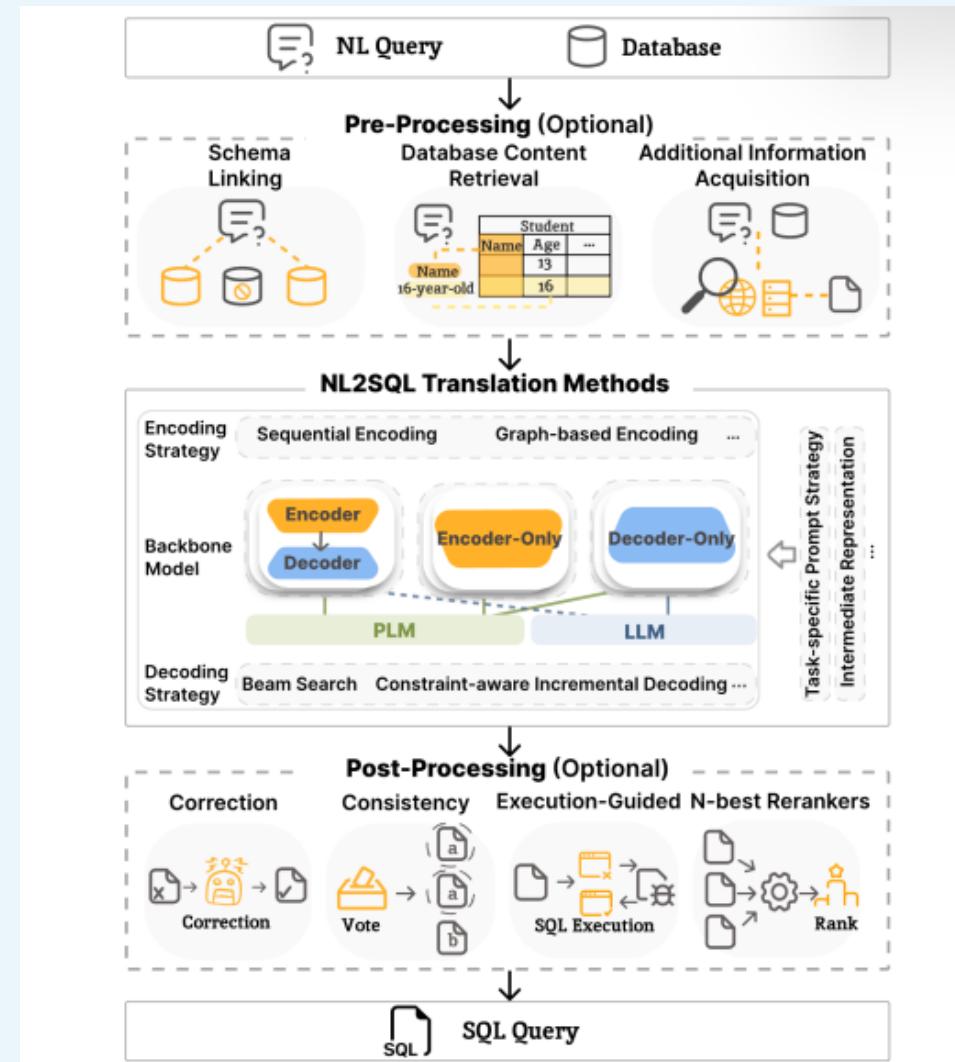


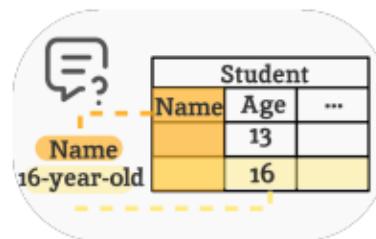
Figure: An Overview of NL2SQL Methods in the LM Era

PRE-PROCESSING

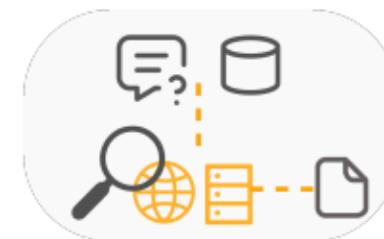
- The pre-processing ensures the accurate mapping and processing of key information within the limited input, by identifying the tables and columns related to the given NL query.
- The pre-processing serves as an enhancement to the model's inputs in the NL2SQL translation.



Schema Linking



Database Content Retrieval



Additional Information Acquisition

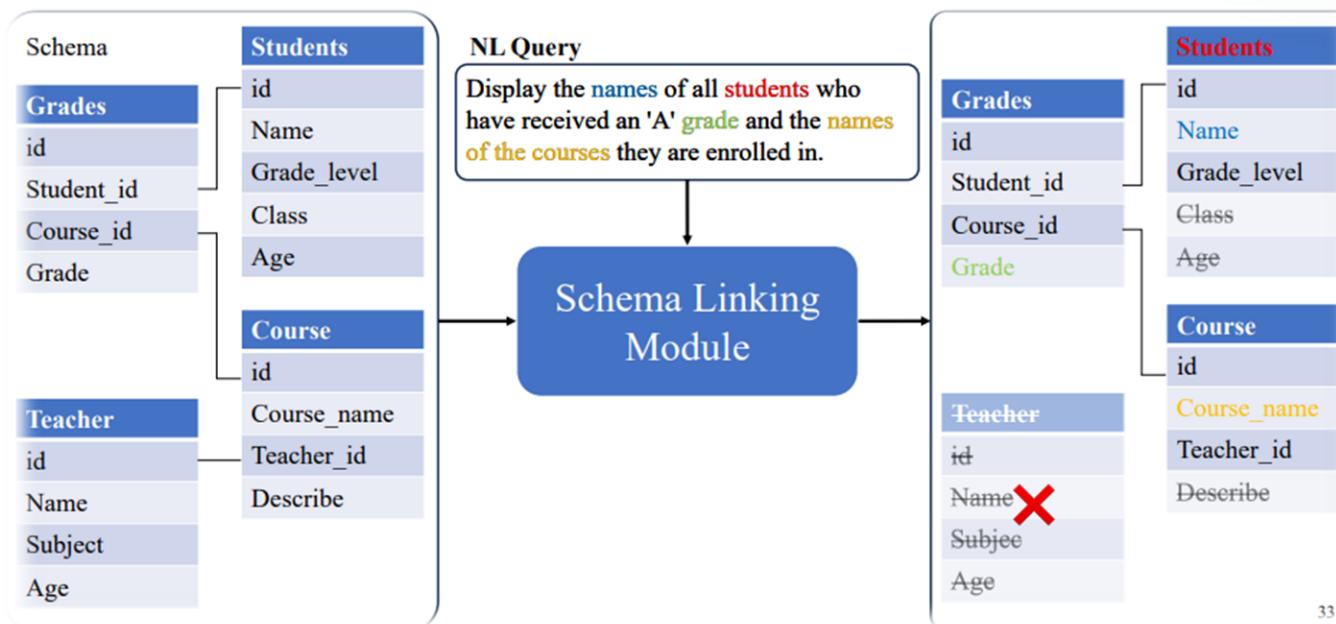
SCHEMA LINKING

Motivation:

- Schema complexity and presence of irrelevant tables with NL in database.
- The model has input limitations, e.g., the input limit of ChatGPT is 4096 tokens.

Goal:

- Identify the tables and columns related to the given NL.



METHOD CLASSIFICATION OF SCHEME LINKING

Method classification :

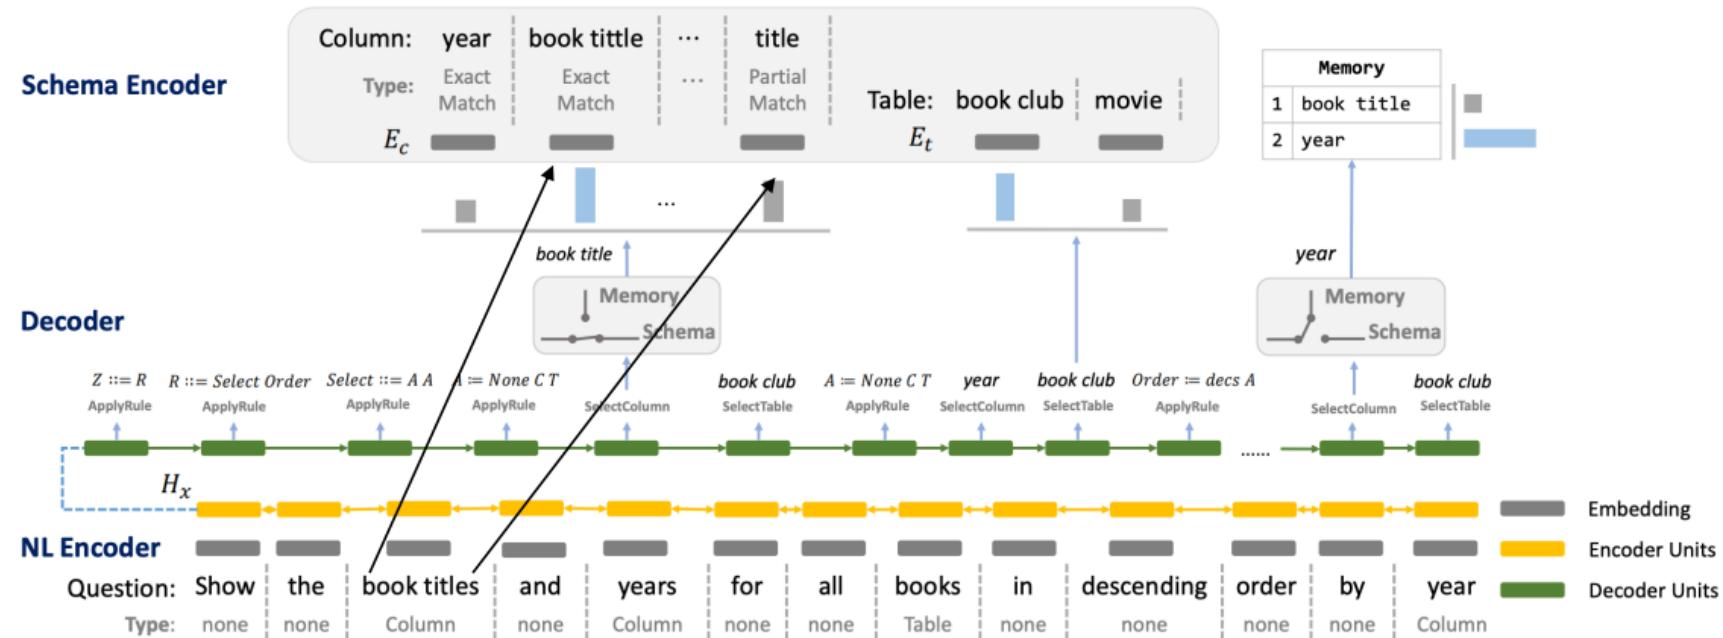
- String Matching-based Schema Linking
- Neural Network-based Schema Linking
- In-Context Learning for Schema Linking



STRING MATCHING-BASED SCHEMA LINKING

IRNet designs string-based schema linking and an intermediate representation called SemQL.

- It uses n-grams (1-6 length) extracted from user queries as query candidates
- It evaluates the string-level similarity between the query candidates and the database schema, then classifies them as Exact or Partial matches.



Jiaqi Guo, Zecheng Zhan, Yan Gao, Yan Xiao, Jian-Guang Lou, Ting Liu, Dongmei Zhang:
Towards Complex Text-to-SQL in Cross-Domain Database with Intermediate Representation. ACL (1) 2019: 4524-4535

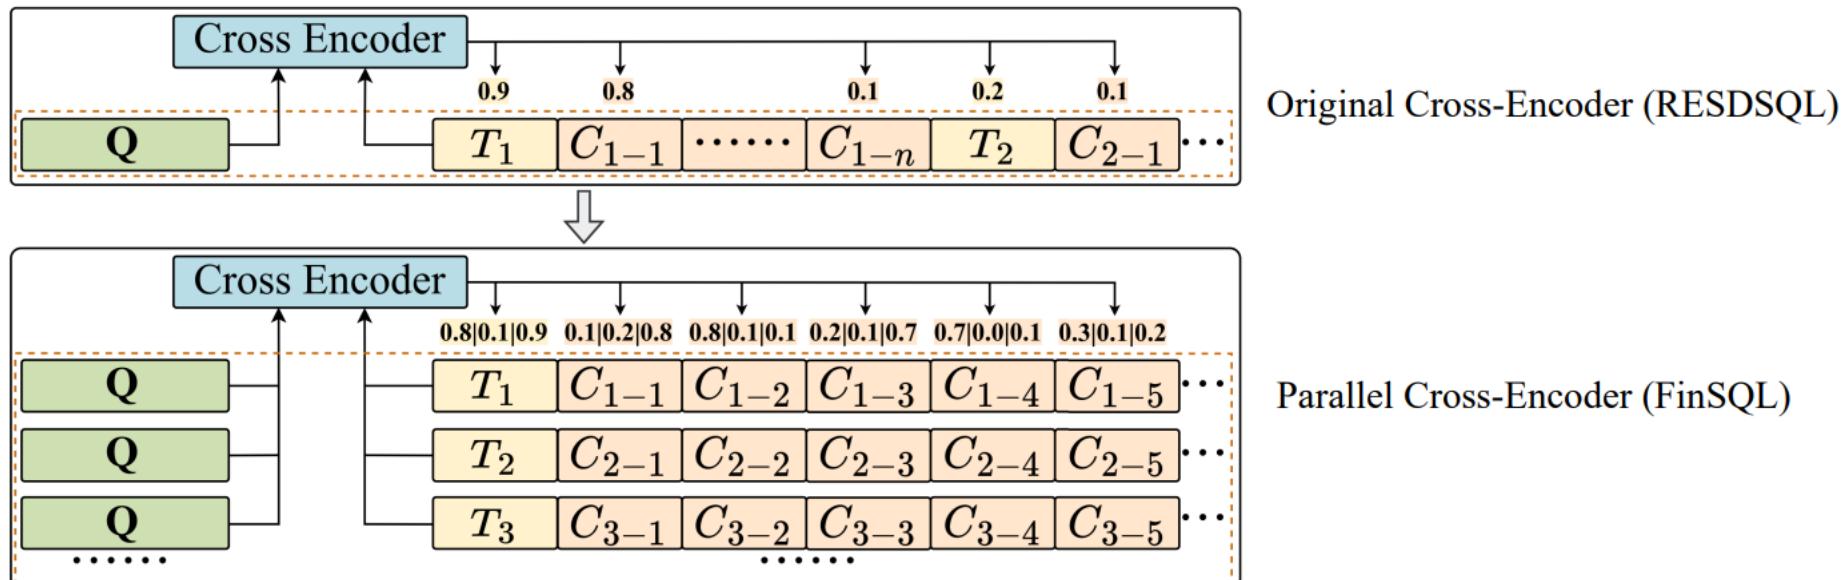
NEURAL NETWORK-BASED SCHEMA LINKING

RESDSQL proposes a ranking-enhanced encoding and skeleton-aware decoding framework.

- A cross-encoder is trained to classify tables and columns based on the input query.

FinSQL, a model-agnostic LLMs-based NL2SQL framework for financial analysis.

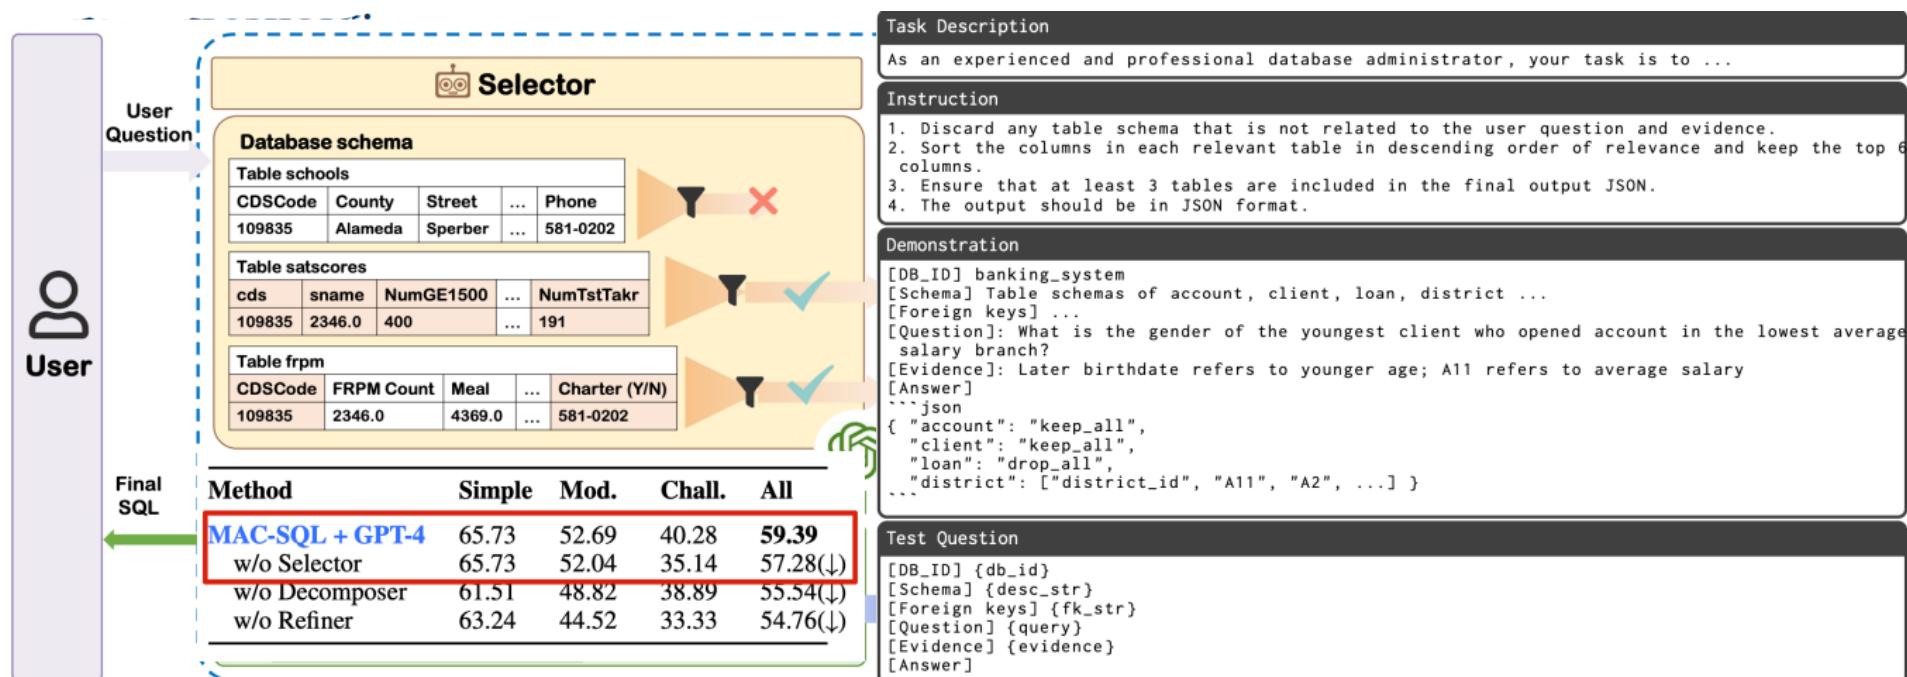
- Organizing the tables into a batch.
- Devising a Parallel Cross-Encoder to retrieve relevant tables and columns from hundreds of schema items



IN-CONTEXT LEARNING FOR SCHEMA LINKING

MAC-SQL is a multi-agent collaborative framework for NL2SQL.

- The Selector agent performs the schema linking task.
- The Selector is activated only when the length of the database schema prompt exceeds a specified threshold



DATABASE CONTENT RETRIEVAL

Motivation:

- Due to the large databases and input limitations of the model, retrieving cell values is resource-intensive.
- Effectively handle some SQL clauses such as WHERE and JOIN .

Goal:

- Efficiently retrieve cell values related to the given NL.

METHOD CLASSIFICATION OF DB CONTENT RETRIEVAL

Method classification :

- String Matching-based Database Content Retrieval
- Neural Network-based Database Content Retrieval
- Index Strategy for Database Content Retrieval

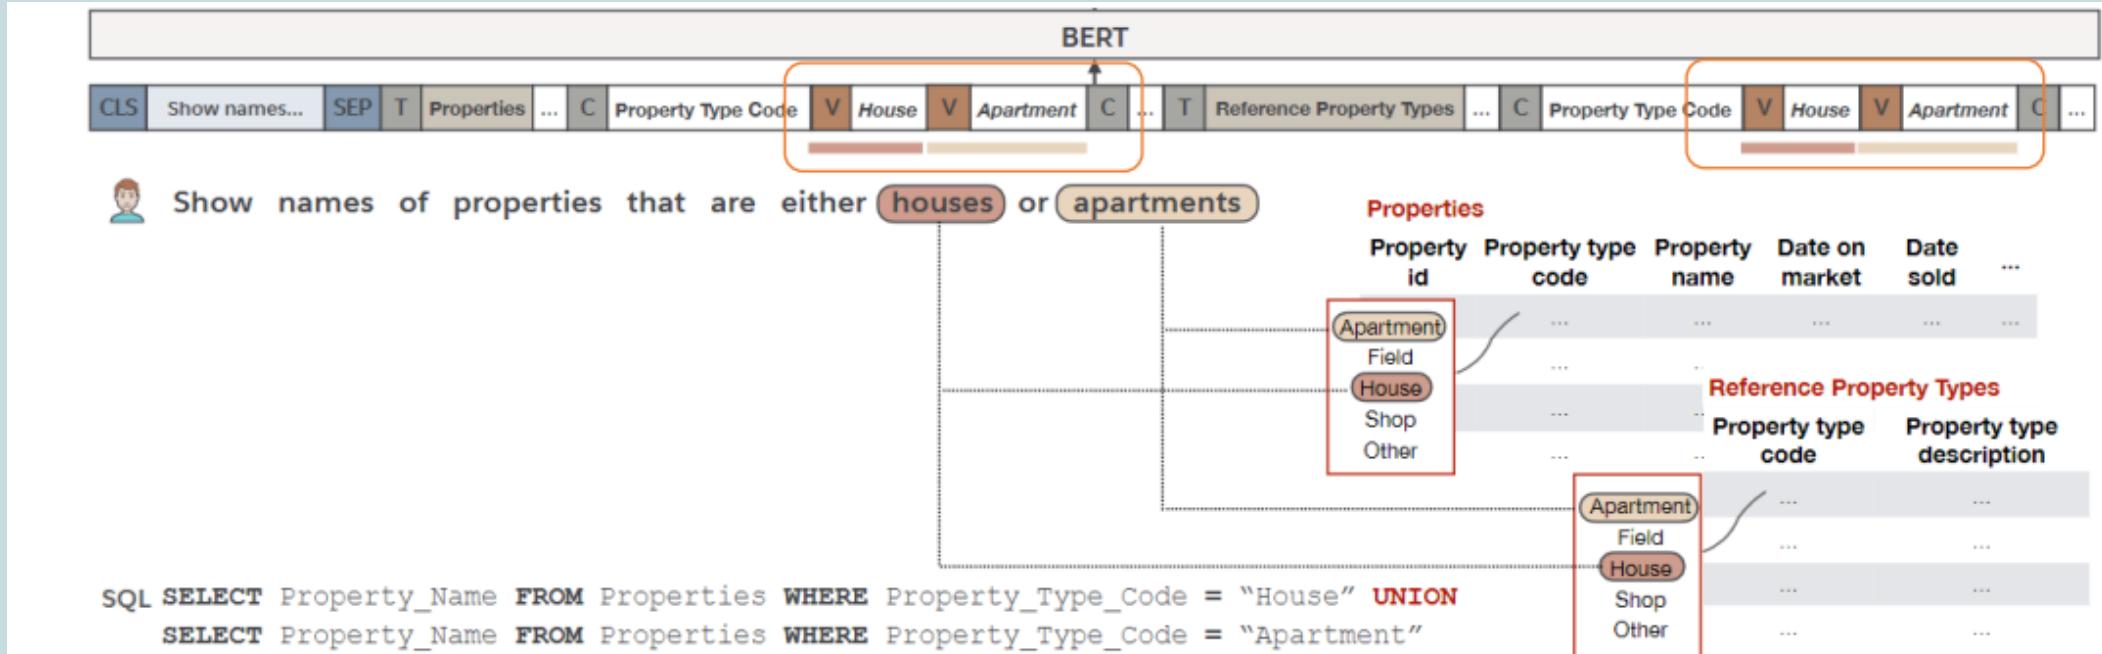
Student		
Name	Age	...
16-year-old	13	
16	16	

Database Content Retrieval

STRING MATCHING-BASED DATABASE CONTENT RETRIEVAL

BRIDGE designs an anchor text matching to automatically extract cell values mentioned in the NL.

- It uses a heuristic method to calculate the maximum sequence match between the problem and the cell values to determine the matching boundary.



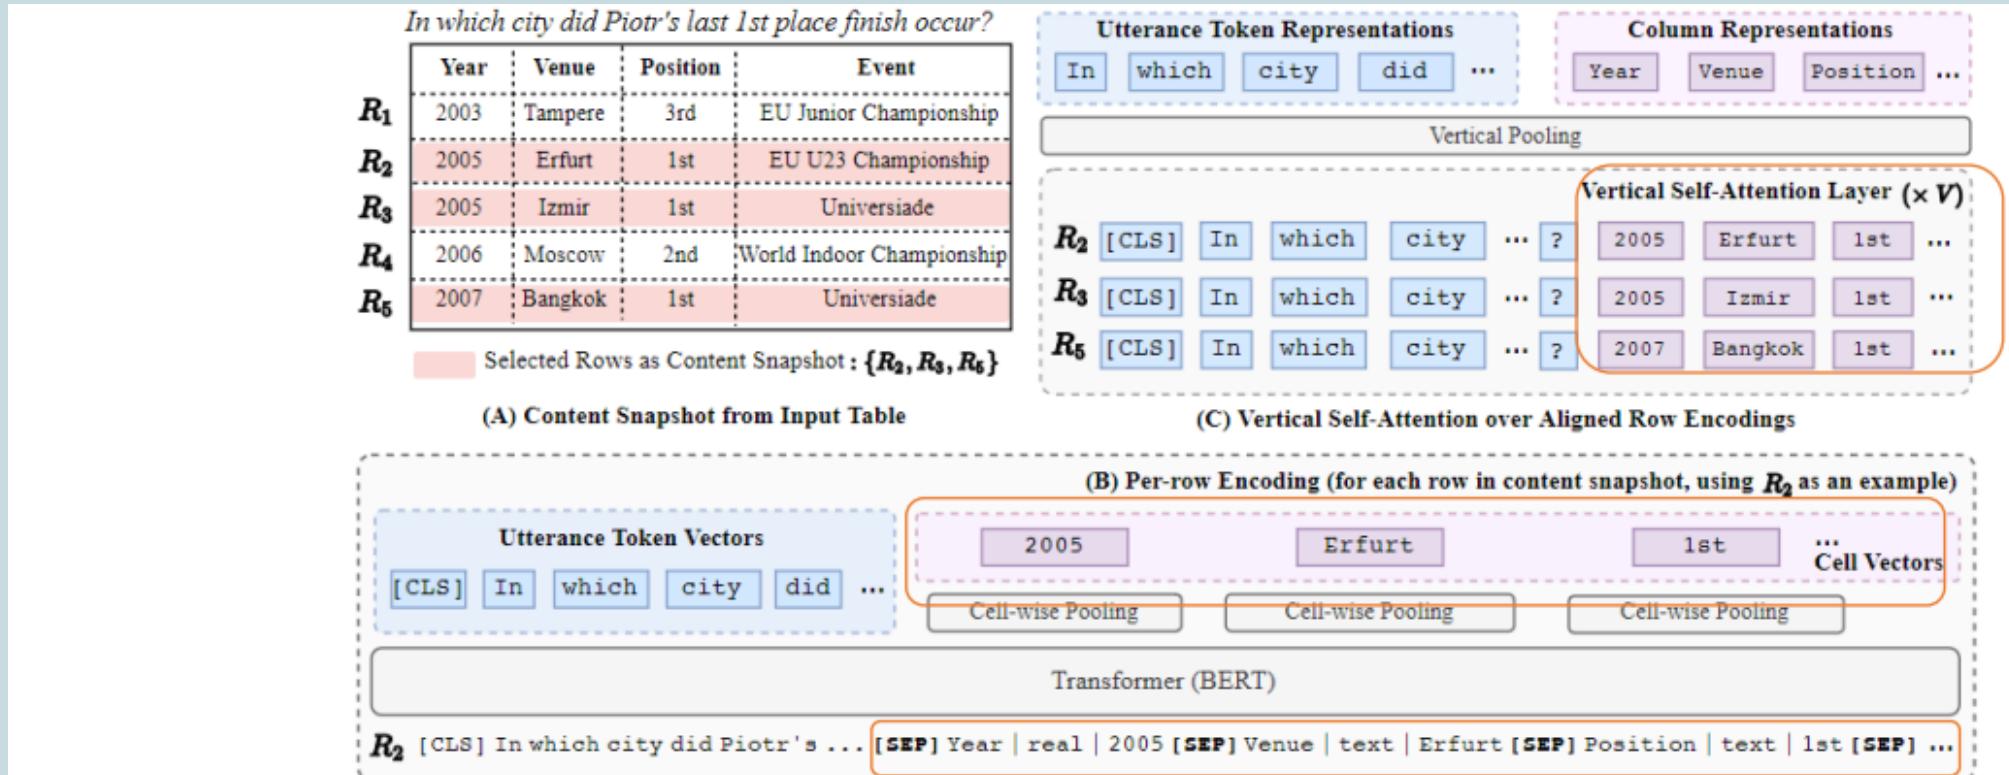
Xi Victoria Lin, Richard Socher, Caiming Xiong:

Bridging Textual and Tabular Data for Cross-Domain Text-to-SQL Semantic Parsing. EMNLP (Findings) 2020: 4870-4888

NEURAL NETWORK-BASED DATABASE CONTENT RETRIEVAL

TaBERT uses a method called database content snapshots.

- It uses an attention mechanism to manage information between cell values across different rows.



Pengcheng Yin, Graham Neubig, Wen-tau Yih, Sebastian Riedel:

TaBERT: Pretraining for Joint Understanding of Textual and Tabular Data. ACL 2020: 8413-8426

NEURAL NETWORK-BASED DATABASE CONTENT RETRIEVAL

Graphix-T5 improves structural reasoning capabilities by modeling the relationship between cell values and the NL query.

- it uses a value-match relation that defines the relationship between cell values and questions.

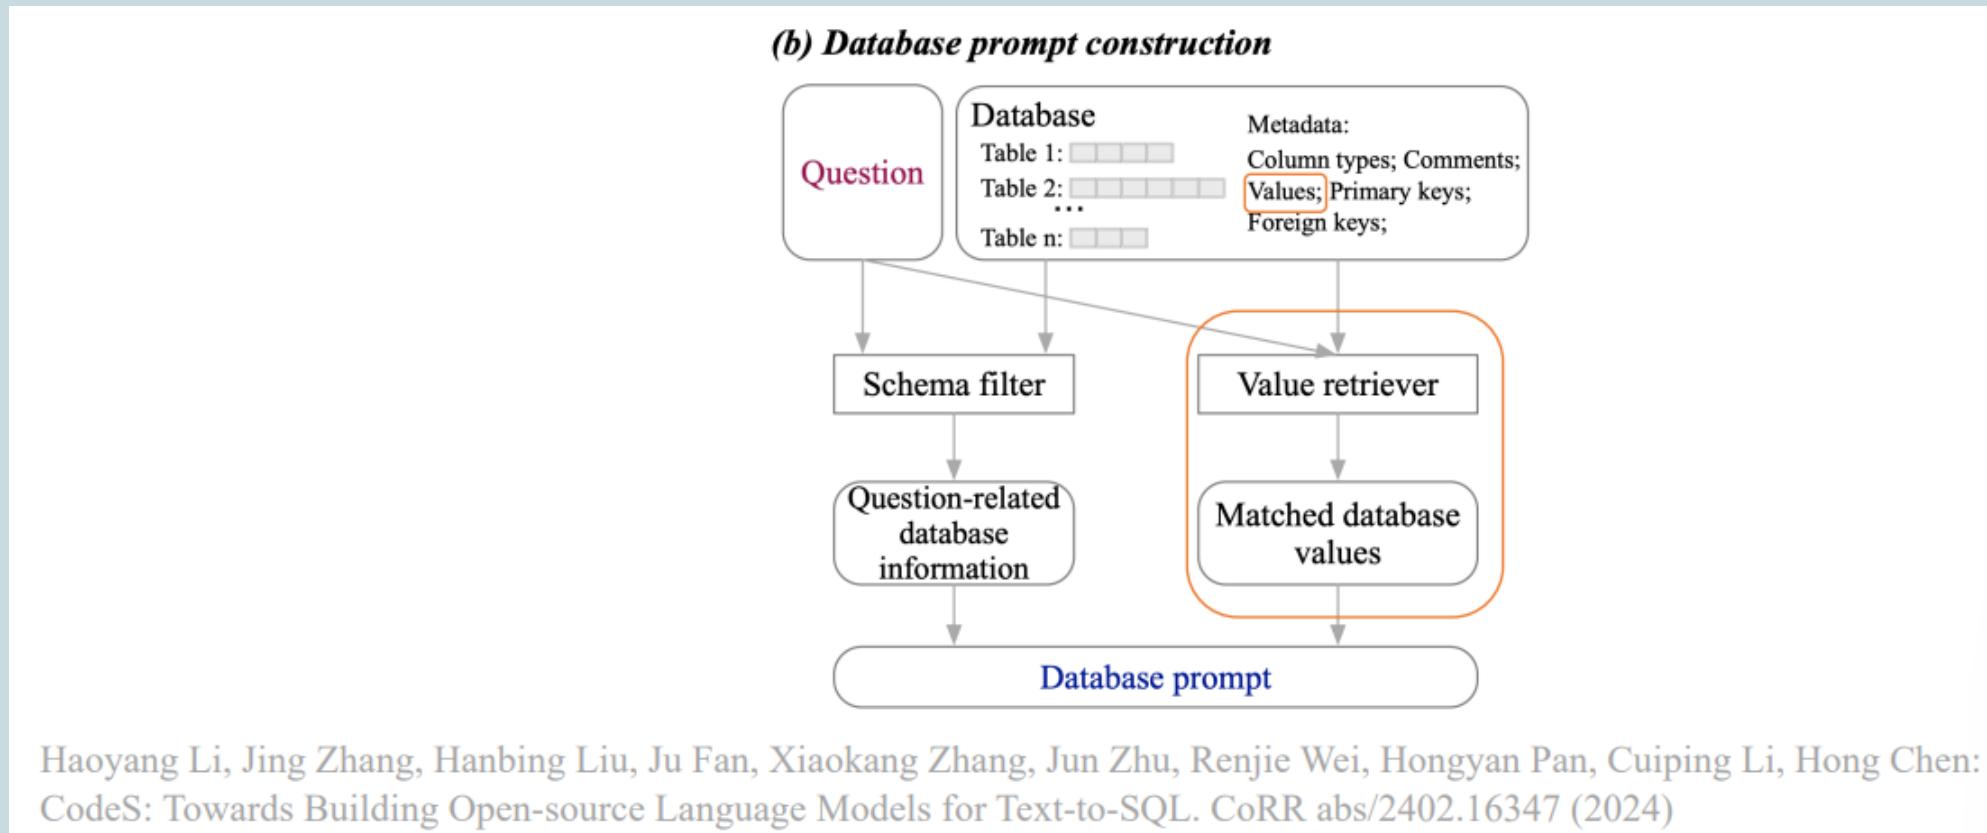


Question	Column	VALUE-MATCH	x is part of the candidate cell values of column y.
----------	--------	-------------	---

INDEX STRATEGY FOR DATABASE CONTENT RETRIEVAL

CodeS introduces a coarse-to-fine cell value matching approach.

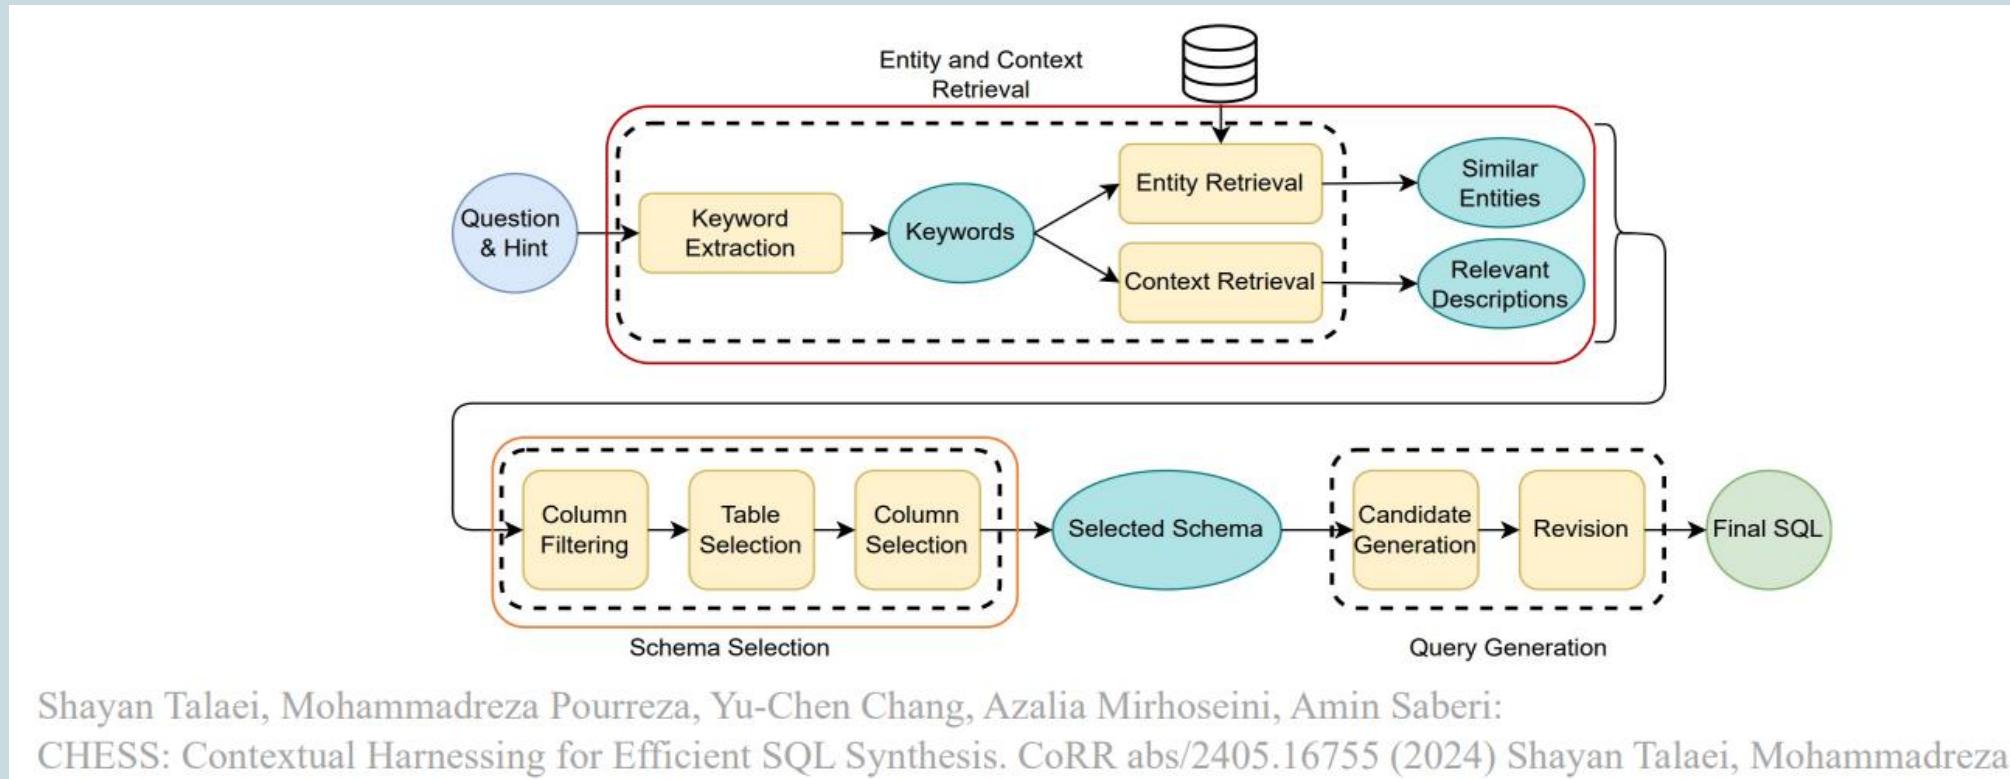
- It builds the index for all values using BM25.
- The index identifies candidate values that are relevant to NL.



INDEX STRATEGY FOR DATABASE CONTENT RETRIEVAL

CHESS uses a Locality-sensitive Hashing algorithm for approximate nearest neighbor searches.

- It indexes unique cell values to quickly identify the top similar values related to the NL query



Shayan Talaei, Mohammadreza Pourreza, Yu-Chen Chang, Azalia Mirhoseini, Amin Saberi:

CHESS: Contextual Harnessing for Efficient SQL Synthesis. CoRR abs/2405.16755 (2024) Shayan Talaei, Mohammadreza

ENCODING STRATEGY

Motivation:

- Converts unstructured and semi-structured data into a form that can be processed for generating SQL queries.
- The encoding process captures the semantic meaning of the NL input and the structural information of the database schema, enabling the model to understand and map the user's intent to the corresponding SQL query.

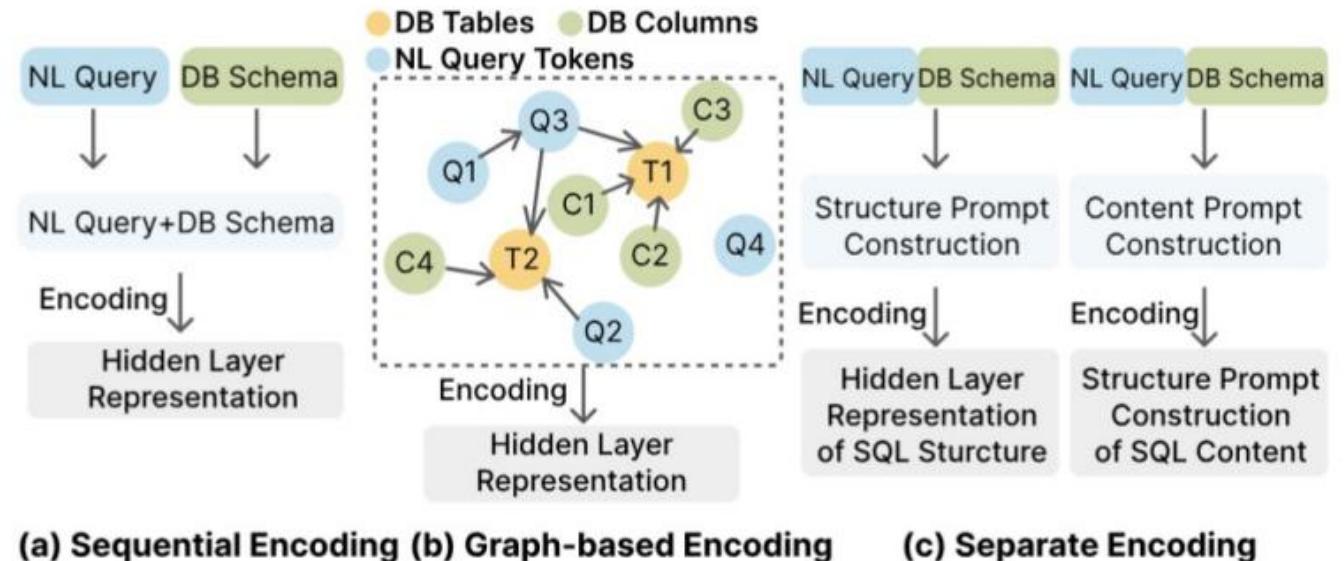
Goal:

- Transform NL and database schema into a structured format that can be effectively utilized by a language model.

METHOD CLASSIFICATION OF ENCODING STRATEGY

Method classification :

- Sequential Encoding Strategy
- Graph-based Encoding Strategy
- Separate Encoding Strategy



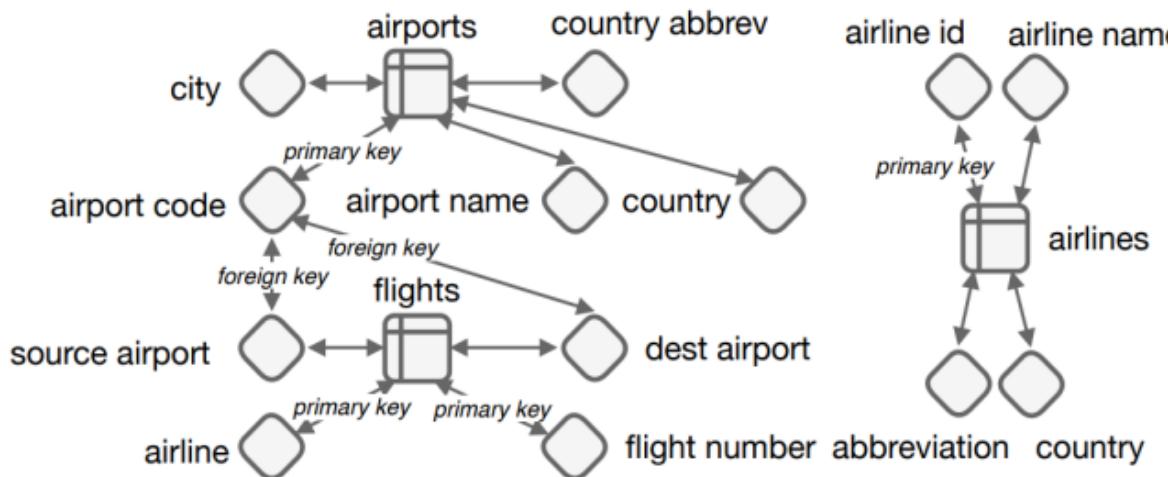
An Overview of the Encoding Strategies

GRAPH-BASED ENCODING STRATEGY

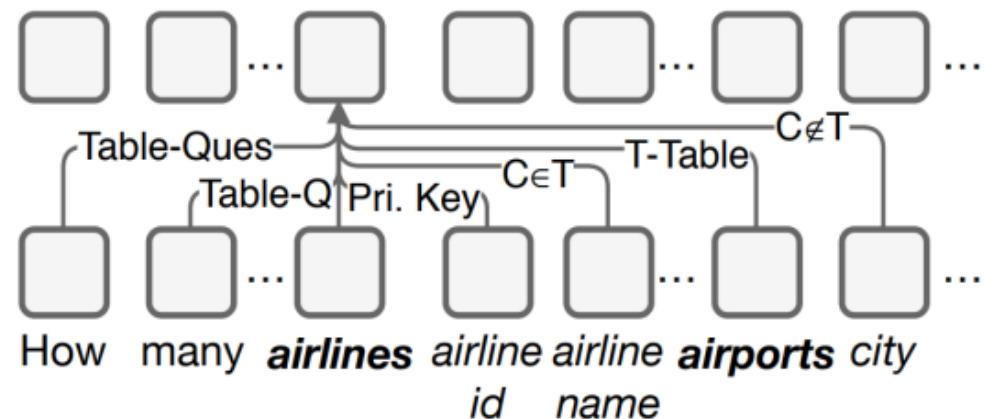
RAT-SQL introduces a relation-aware self-attention mechanism.

- It allows the model to explicitly consider and utilize predefined relational information when jointly encoding the question and the database schema.
- These relationships are represented as a graph structure, and it can more effectively capture the structural information in the schema and its alignment with the NL query

An illustration of an example schema as a graph G



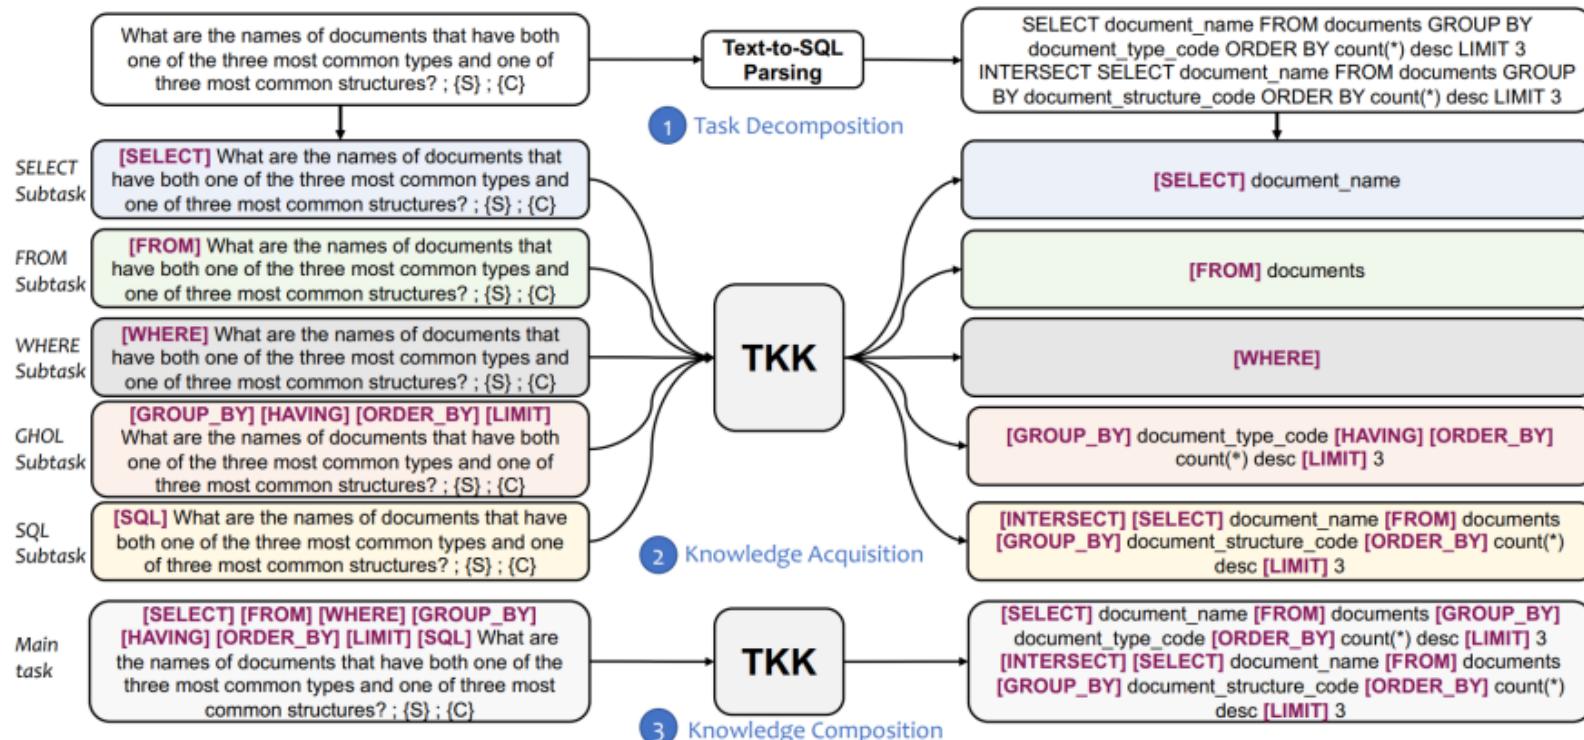
One RAT layer in the schema encoder.



SEPARATE ENCODING STRATEGY

TKK employs task decomposition and multi-task learning strategies in encoding.

- It breaks down the complex NL2SQL task into multiple subtasks.
- In this way, it can progressively acquire and combine knowledge.



DECODING STRATEGY

Motivation:

- The choice of decoding strategy directly affects the quality and performance of the generated SQL queries.
- An excellent decoding strategy not only produces syntactically correct SQL queries but also ensures that the semantics of the SQL queries align with the NL and can even optimize the execution efficiency of the queries.

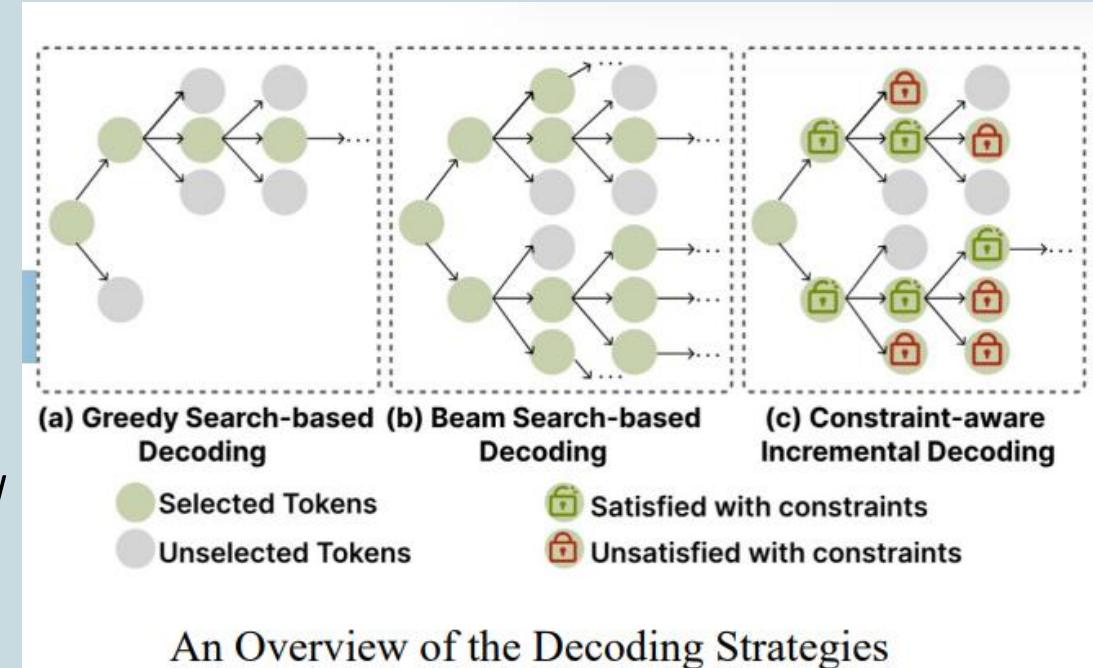
Goal:

- Convert the representations generated by the encoder into the target SQL queries.

METHOD CLASSIFICATION OF DECODING STRATEGY

Method classification :

- Greedy search-based decoding strategy
- Beam search-based decoding strategy
- Constraint-aware incremental decoding strategy



GREEDY SEARCH-BASED DECODING STRATEGY

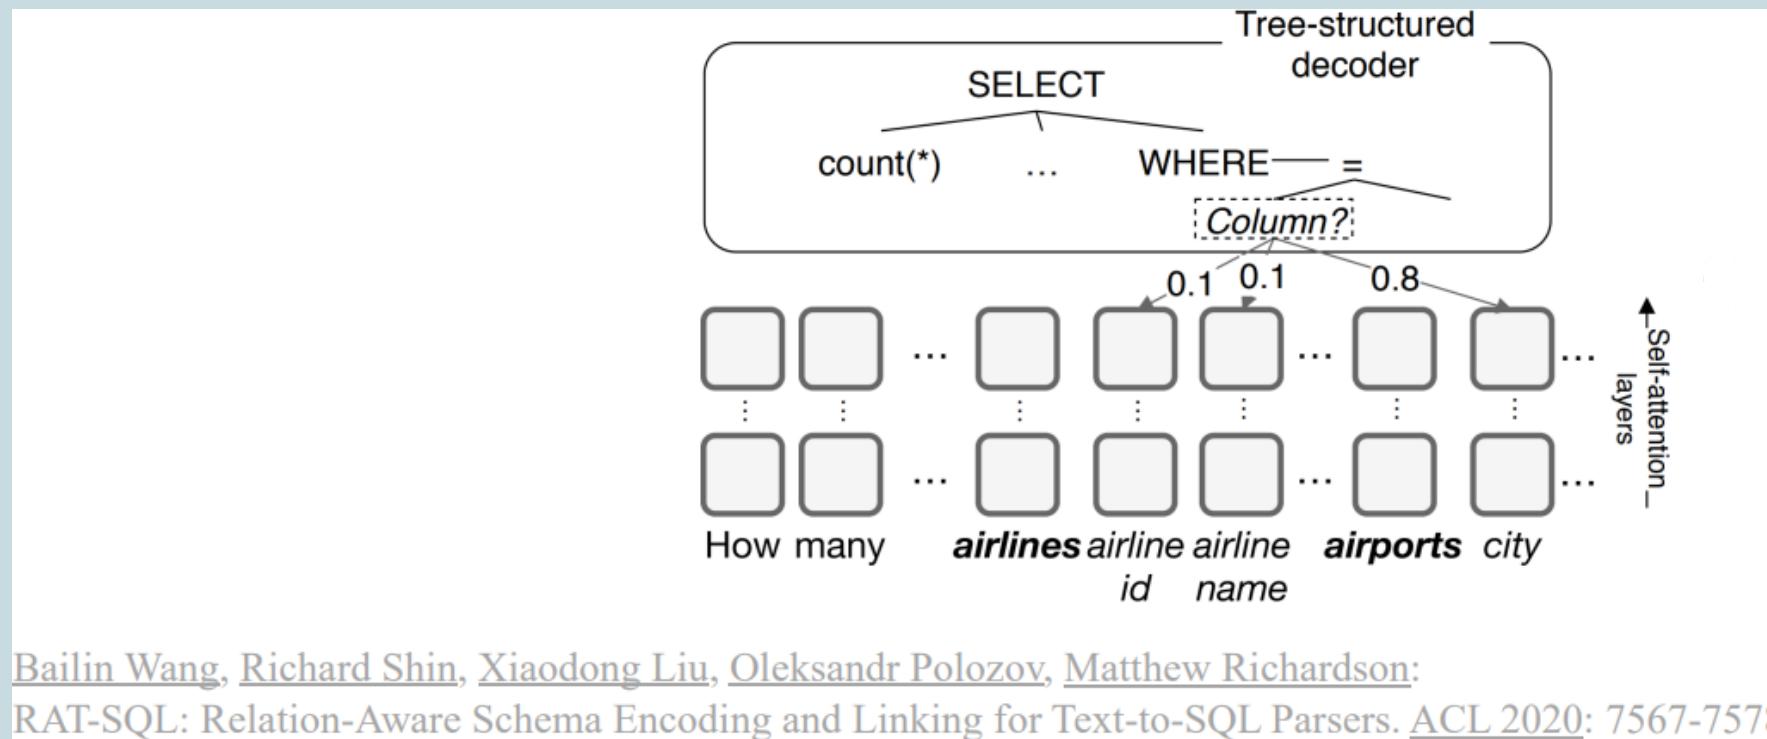
The greedy search-based decoding strategy is a simple and fast approach for decoding.

- At each decoding step, greedy search selects the token with the highest current probability as the output.
- This strategy builds the final output sequence by continuously choosing the locally optimal solution.
- Since the default decoding strategy of GPT series models (e.g., GPT-4) is greedy search-based decoding, NL2SQL solutions based on GPT fall into this category.

BEAM SEARCH-BASED DECODING STRATEGY

RAT-SQL combines relation-aware graph structure encoding and generation techniques.

- During the decoding process, RAT-SQL uses beam search to generate multiple candidate SQL queries
 - These queries are then reranked, and the optimal query is selected based on graph structure information



CONSTRAINT-AWARE INCREMENTAL DECODING STRATEGY

PICARD (Parsing Incrementally for Constrained Auto-Regressive Decoding) introduces the constraint-aware incremental decoding strategy.

- The constraint-aware incremental decoding strategy is specifically designed for NL2SQL tasks.
- This strategy aims to ensure the generation of syntactically correct SQL queries by incorporating constraints during the decoding process.

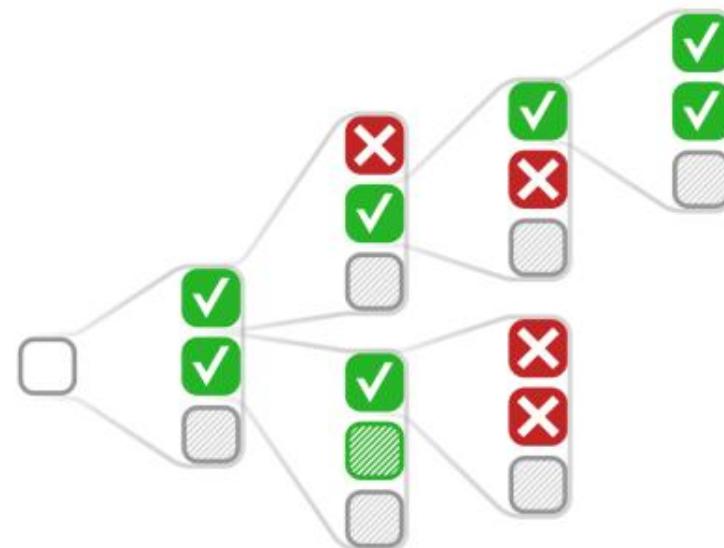


Illustration of constrained beam search with beam size 2 and PICARD.

Torsten Scholak, Nathan Schucher, Dzmitry Bahdanau:

PICARD: Parsing Incrementally for Constrained Auto-Regressive Decoding from Language Models. *EMNLP (1)* 2021: 9895-9901

INTERMEDIATE REPRESENTATION

Motivation:

- Design a grammar-free intermediate representation compared to SQL as the bridge between the "free-form" NL query and the "constrained and formal" SQL query.

Goal:

- Capture the essential components and relationships of an NL query without the strict syntax rules of SQL.



NL Query:

Which film has more than 5 actors and less than 3 in the inventory?

film		
title	film_id	...

film_actor		
name	film_id	...

inventory		
tag	film_id	...

• - • Foreign Key

Intermediate Representation:

SQL-like Syntax Language
(e.g. NATSQL)

```
SELECT film.title  
WHERE count(film_actor.*)>5 And count(inventory.*)<3
```

SQL-like Sketch Structure
(e.g. SC-Prompt)

```
SELECT [column]          [column] title  
FROM [table]           [table] film  
JOIN [table]          [table] film_actor  
ON [table].[column]    [table].[column] film.film_id  
=[table].[column]      [table].[column] film_actor.film_id  
GROUP BY [column]     [column] film_id  
HAVING count([column]) > n [column] * [n] 5  
INTERSECT  
(SELECT [column]          [column] title  
FROM [table]           [table] film  
JOIN [table]          [table] inventory  
ON [table].[column]    [table].[column] film.film_id  
=[table].[column]      [table].[column] inventory.film_id  
GROUP BY [column]     [column] film_id  
HAVING count([column]) < n [column] * [n] 3
```

SQL:

```
SELECT T1.title  
FROM film AS T1 JOIN film_actor AS T2 ON T1.film_id = T2.film_id  
GROUP BY T1.film_id HAVING count(*) > 5  
INTERSECT  
SELECT T1.title  
FROM film AS T1 JOIN inventory AS T2 ON T1.film_id = T2.film_id  
GROUP BY T1.film_id HAVING count(*) < 3
```

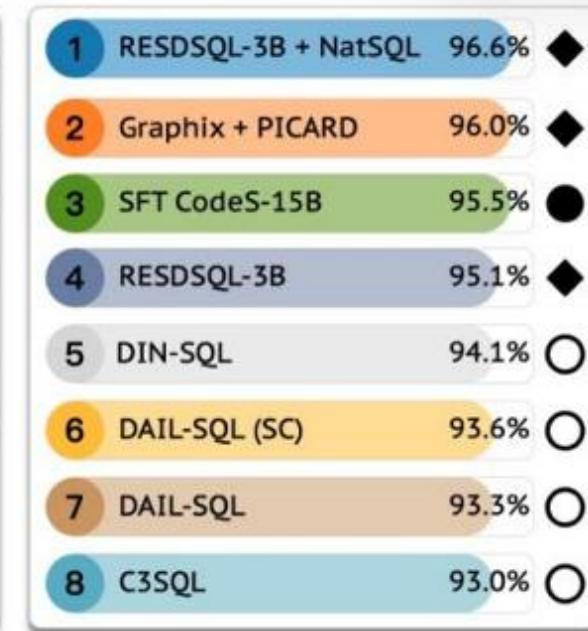
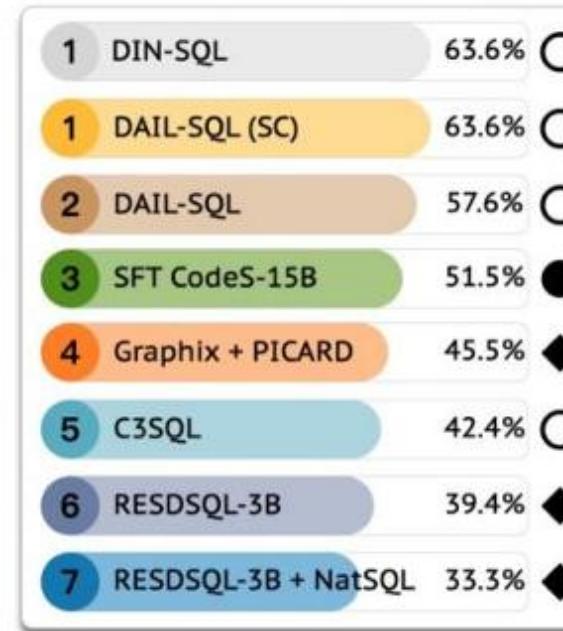
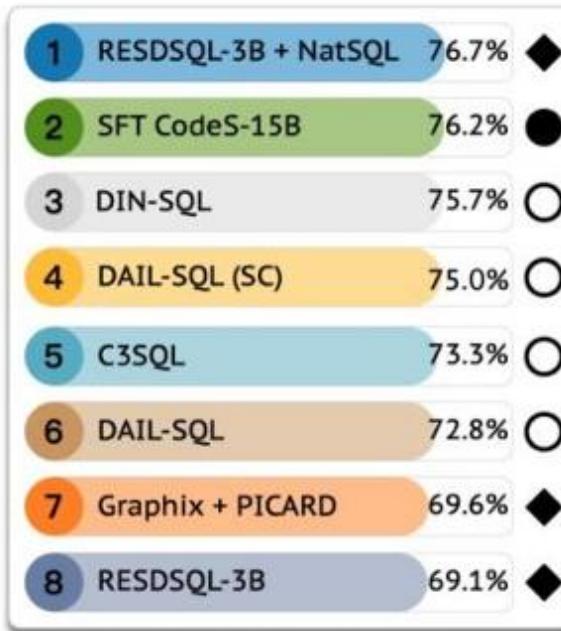
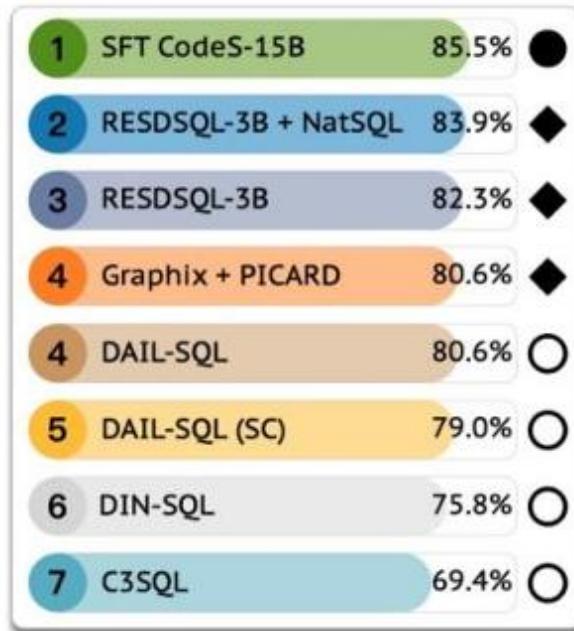
POST PROCESSING

Motivation:

- Post-processing is a crucial step to refine the generated SQL queries, ensuring they meet user expectations more accurately.
- This involves enhancing the initial SQL output using various strategies.



NL2SQL360: LEADERBOARD



(a) Test Results in Domain of Competition

(b) Test Results in SQL with JOIN

(c) Test Results in Nested SQL

(d) Test Results in Query Variance

Figure 3: NL2SQL Models on Spider from Different Angles (○: Prompting LLM, ●: Fine-tuning LLM, ◆: Fine-tuning PLM).

REFERENCES

- Liu, X., Shen, S., Li, B., Ma, P., Jiang, R., Zhang, Y., Fan, J., Li, G., Tang, N., & Luo, Y. (2024). *A Survey of NL2SQL with Large Language Models: Where are we, and where are we going?*
- Jiaqi Guo, Zecheng Zhan, Yan Gao, Yan Xiao, Jian-Guang Lou, Ting Liu, Dongmei Zhang:Towards Complex Text-to-SQL in Cross-Domain Database with Intermediate Representation. ACL (1) 2019: 4524-4535
- Chao Zhang, Yuren Mao, Yijiang Fan, Yu Mi, Yunjun Gao, Lu Chen, Dongfang Lou, Jinshu Lin:FinSQL: Model-Agnostic LLMs-based Text-to-SQL Framework for Financial Analysis. SIGMOD Conference Companion 2024: 93-105
- Shayan Talaei, Mohammadreza Pourreza, Yu-Chen Chang, Azalia Mirhoseini, Amin Saberi:CHESS: Contextual Harnessing for Efficient SQL Synthesis. CoRR abs/2405.16755 (2024) Shayan Talaei, Mohammadreza
- Bing Wang, Changyu Ren, Jian Yang, Xinnian Liang, Jiaqi Bai, Qian-Wen Zhang, Zhao Yan, Zhoujun Li:MAC-SQL: A Multi-Agent Collaborative Framework for Text-to-SQL. CoRR abs/2312.11242 (2023)
- Pengcheng Yin, Graham Neubig, Wen-tau Yih, Sebastian Riedel-TaBERT: Pretraining for Joint Understanding of Textual and Tabular Data. ACL 2020: 8413-8426
- Haoyang Li, Jing Zhang, Hanbing Liu, Ju Fan, Xiaokang Zhang, Jun Zhu, Renjie Wei, Hongyan Pan, Cuiping Li, Hong Chen:CodeS: Towards Building Open-source Language Models for Text-to-SQL. CoRR abs/2402.16347 (2024)
- Chang Gao, Bowen Li, Wenzuan Zhang, Wai Lam, Binhu Li, Fei Huang, Luo Si, Yongbin Li:Towards Generalizable and Robust Text-to-SQL Parsing. EMNLP (Findings) 2022: 2113-2125
- Torsten Scholak, Nathan Schucher, Dzmitry Bahdanau:PICARD: Parsing Incrementally for Constrained Auto-Regressive Decoding from Language Models. EMNLP (1) 2021: 9895-9901

THANK YOU

